

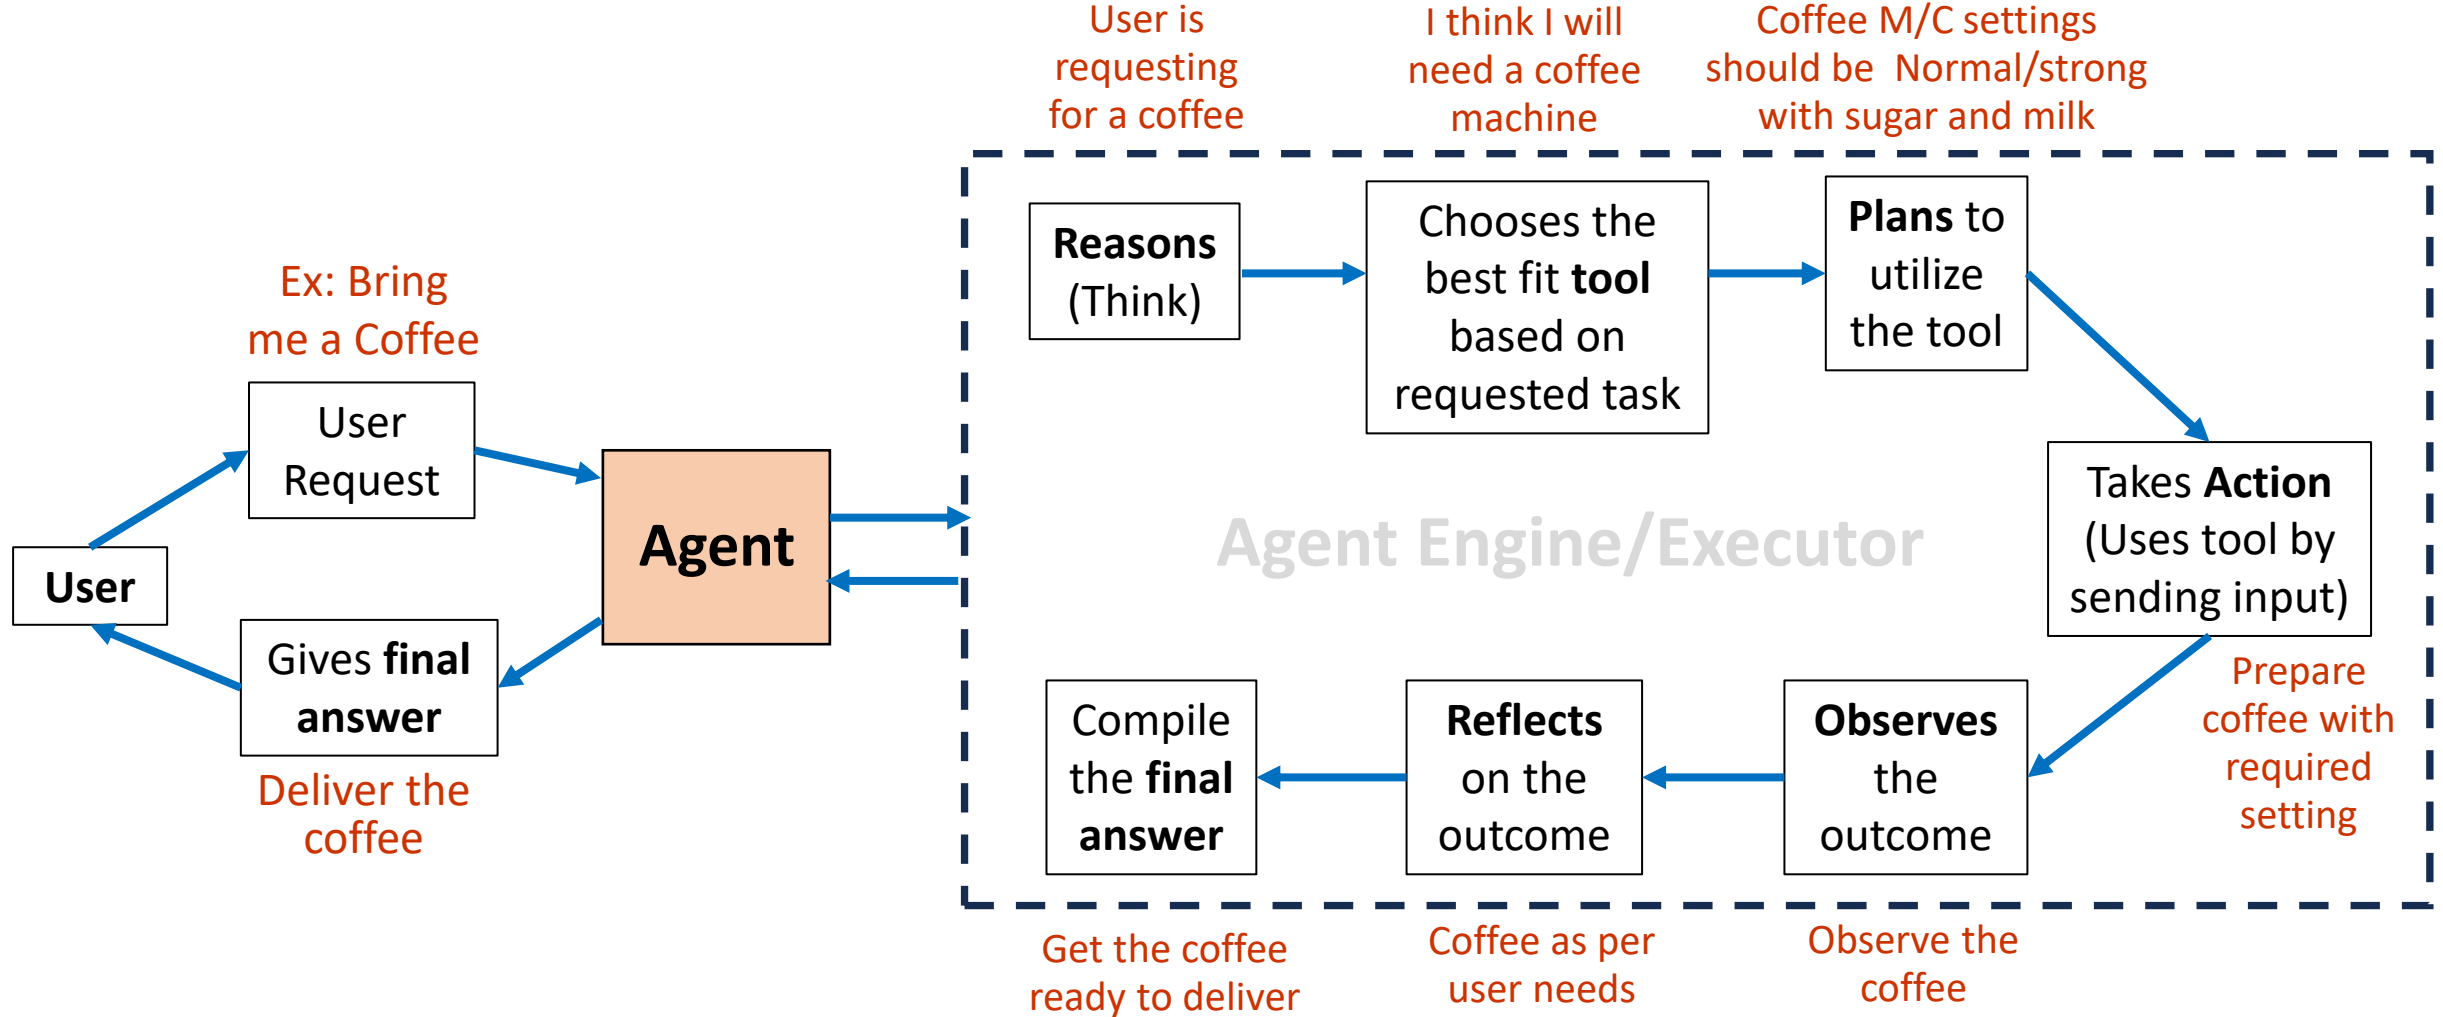


Agents & Agentic Framework

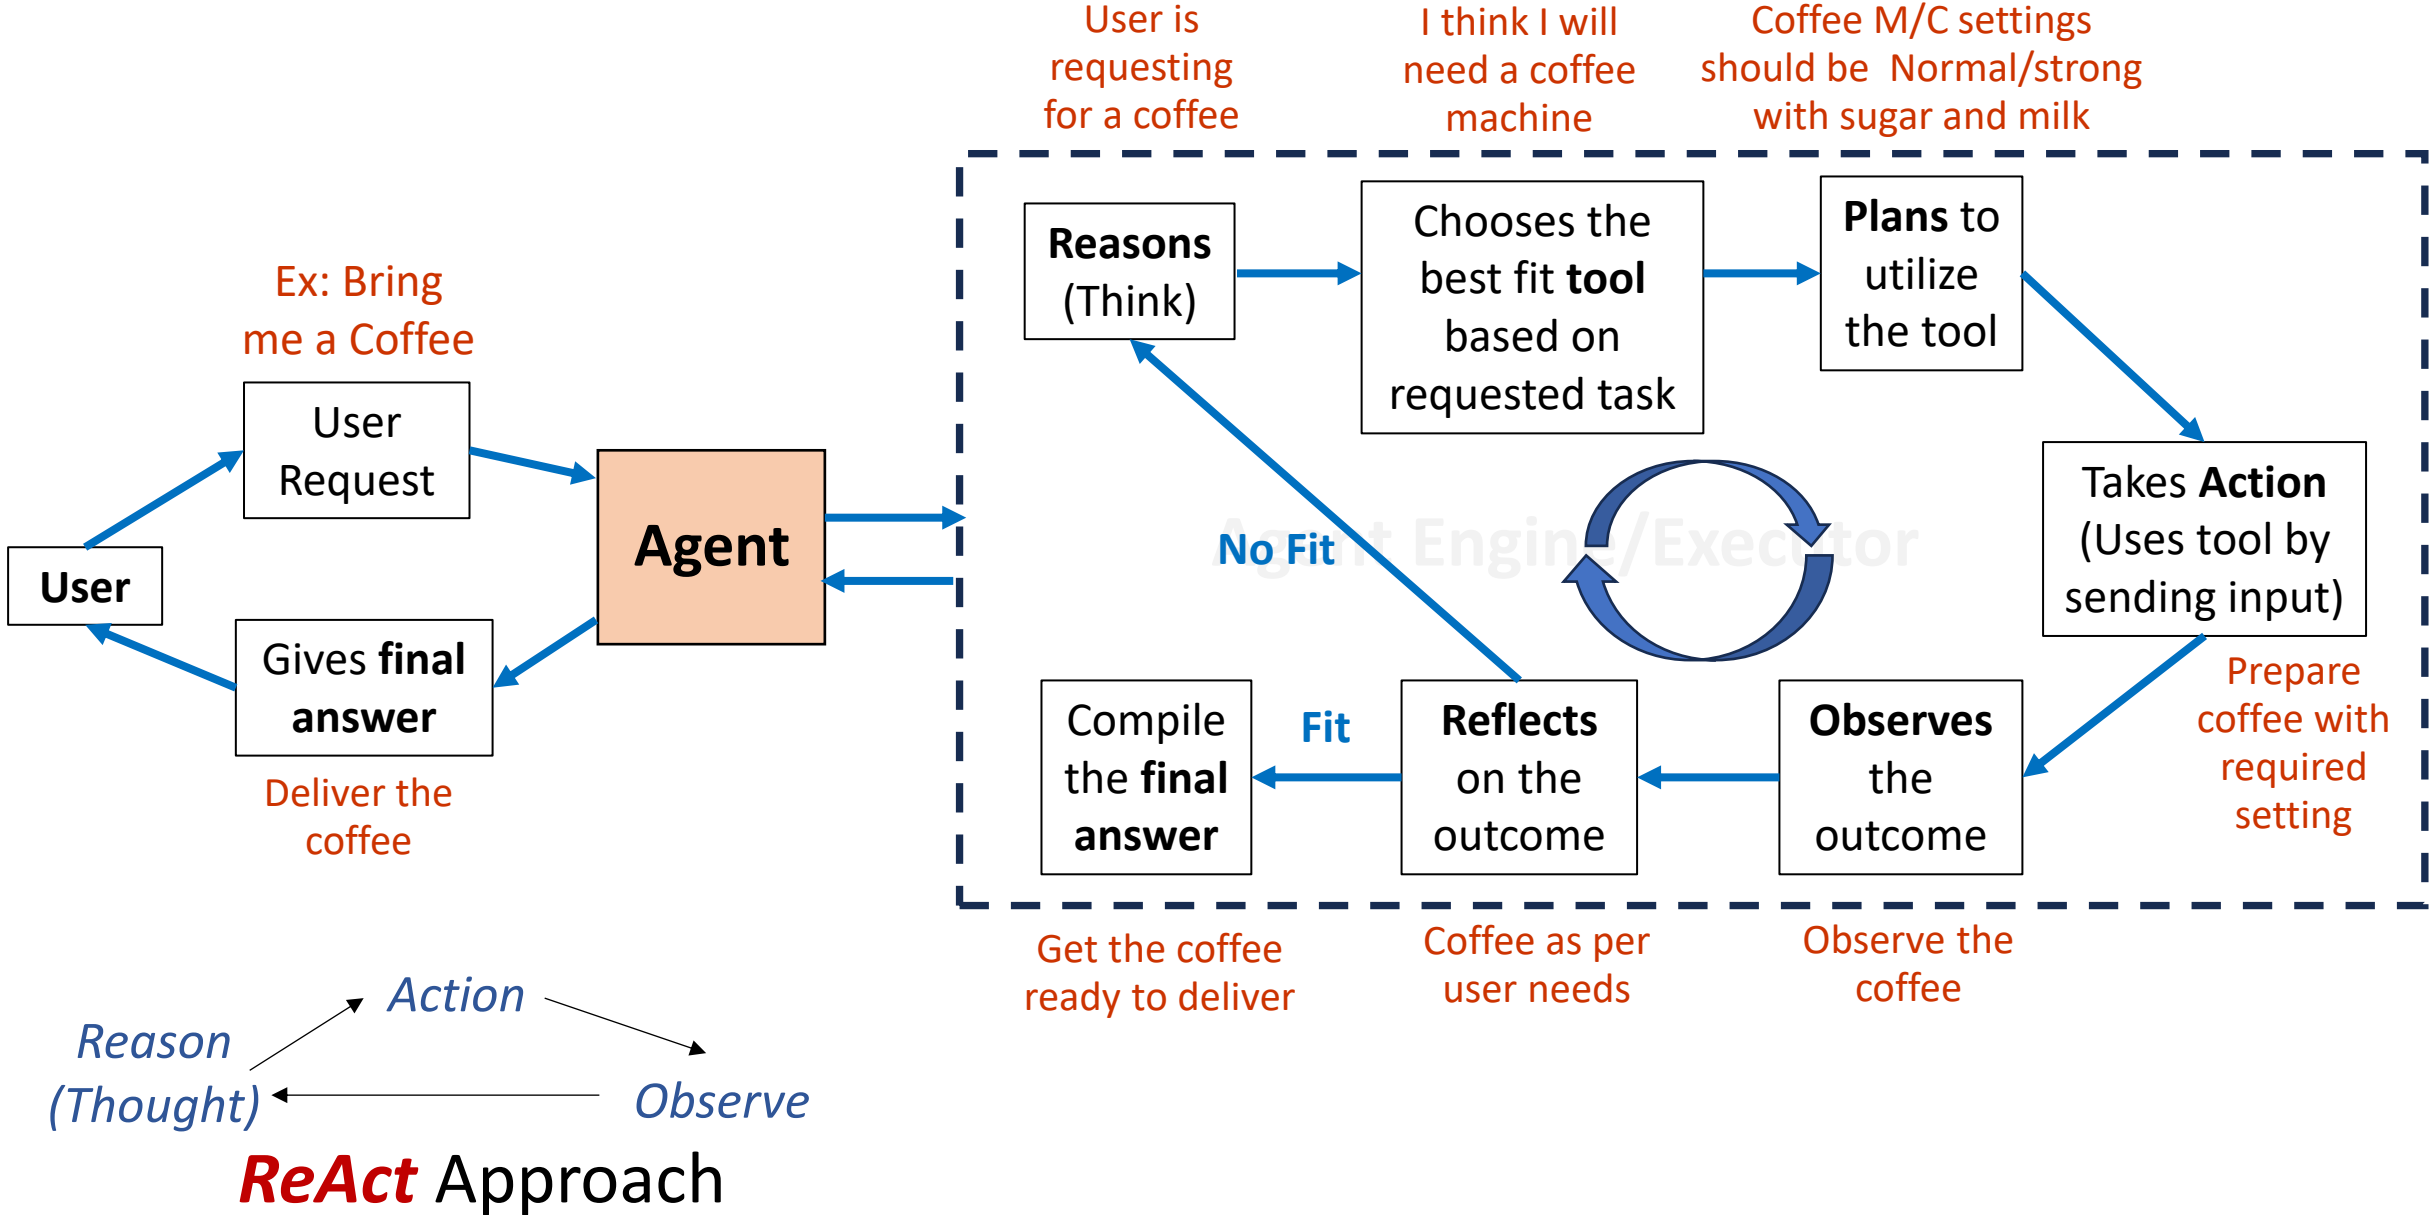
Bhanu Chander V

24th June 2025

Agent - Typical Process



Agentic Framework *is Iterative*



Agent Levels

Agency Level	Description	What that's called
☆☆☆	Agent output has no impact on program flow	Simple processor
★★☆	Agent output determines basic control flow	Router
★★★	Agent output determines function execution	Tool caller

★★★	Agent output controls iteration and program continuation	Multi-step Agent
★★★	One agentic workflow can start another agentic workflow	Multi-Agent

Agentic frameworks involve models that can plan, reason, and act in iterative loops—

- they read a task,
- think step by step, and
- take actions


The Brain — LLMs (for reason & Plan)
Actions — using Tools

- Action – can involve multiple Tools

Examples:

Hey Siri ? Alexa, Google Assistant
Customer Service - Chatbots
Auto Emails

Where MCP is involved (Recap from MCP Training)



Agents Demo

AI Agent Observability and Evaluation

1. **Observability** -> understanding what's happening inside your AI agent by looking at external signals
 - like logs, metrics, and traces.
2. For AI agents, this means tracking actions, tool usage, model calls, and responses to debug and improve agent performance.
3. Observability enables:
 - Understand **costs** and accuracy trade-offs
 - Measure **latency**
 - Detect harmful language & prompt injection
 - Monitor user feedback

Some Tools:

<https://langfuse.com/>

Automated Evaluation Metrics:

<https://docs.ragas.io/en/stable/> (For evaluation of LLM application)

AI Agent Observability and Evaluation

2. Evaluation

Common Metrics to Track in Production

1.Costs — Additional cost for multiple iterations

2.Latency — Observe the time it takes to complete each step, or the entire run.

3.User Feedback — Users can provide direct feedback (thumbs up/down) to help refine or correct the agent.

4.LLM-as-a-Judge — Use a separate LLM to evaluate your agent's output in near real-time (e.g., checking for toxicity or correctness).

Automated Evaluation Metrics:

RAGAS <https://docs.ragas.io/en/stable/> (For evaluation of LLM application)

Additional References/Tools from Langchain:

1. Traces with Langchain - https://docs.smith.langchain.com/observability/how_to_guides/trace_with_langchain
2. Langsmith Tutorial on Observability - <https://docs.smith.langchain.com/>
3. Callbacks can also be used

Human Feedback

You can integrate human feedback at various levels:


Feedback Level	Example Use	Implementation Approach
Intermediate step approval	Before or after tool usage	Custom loop or callbacks
Final answer approval	Let human confirm/modify output	UI, Streamlit, or callback
Training/Scoring data	Evaluate correctness over time	Use LangSmith / logs / metrics
Retrain agent behavior	Reward/punish model decisions	RLHF (advanced, post-logging)

An Opensource Library for Agents:

Smolagents: <https://huggingface.co/docs/smolagents/en/index>

END

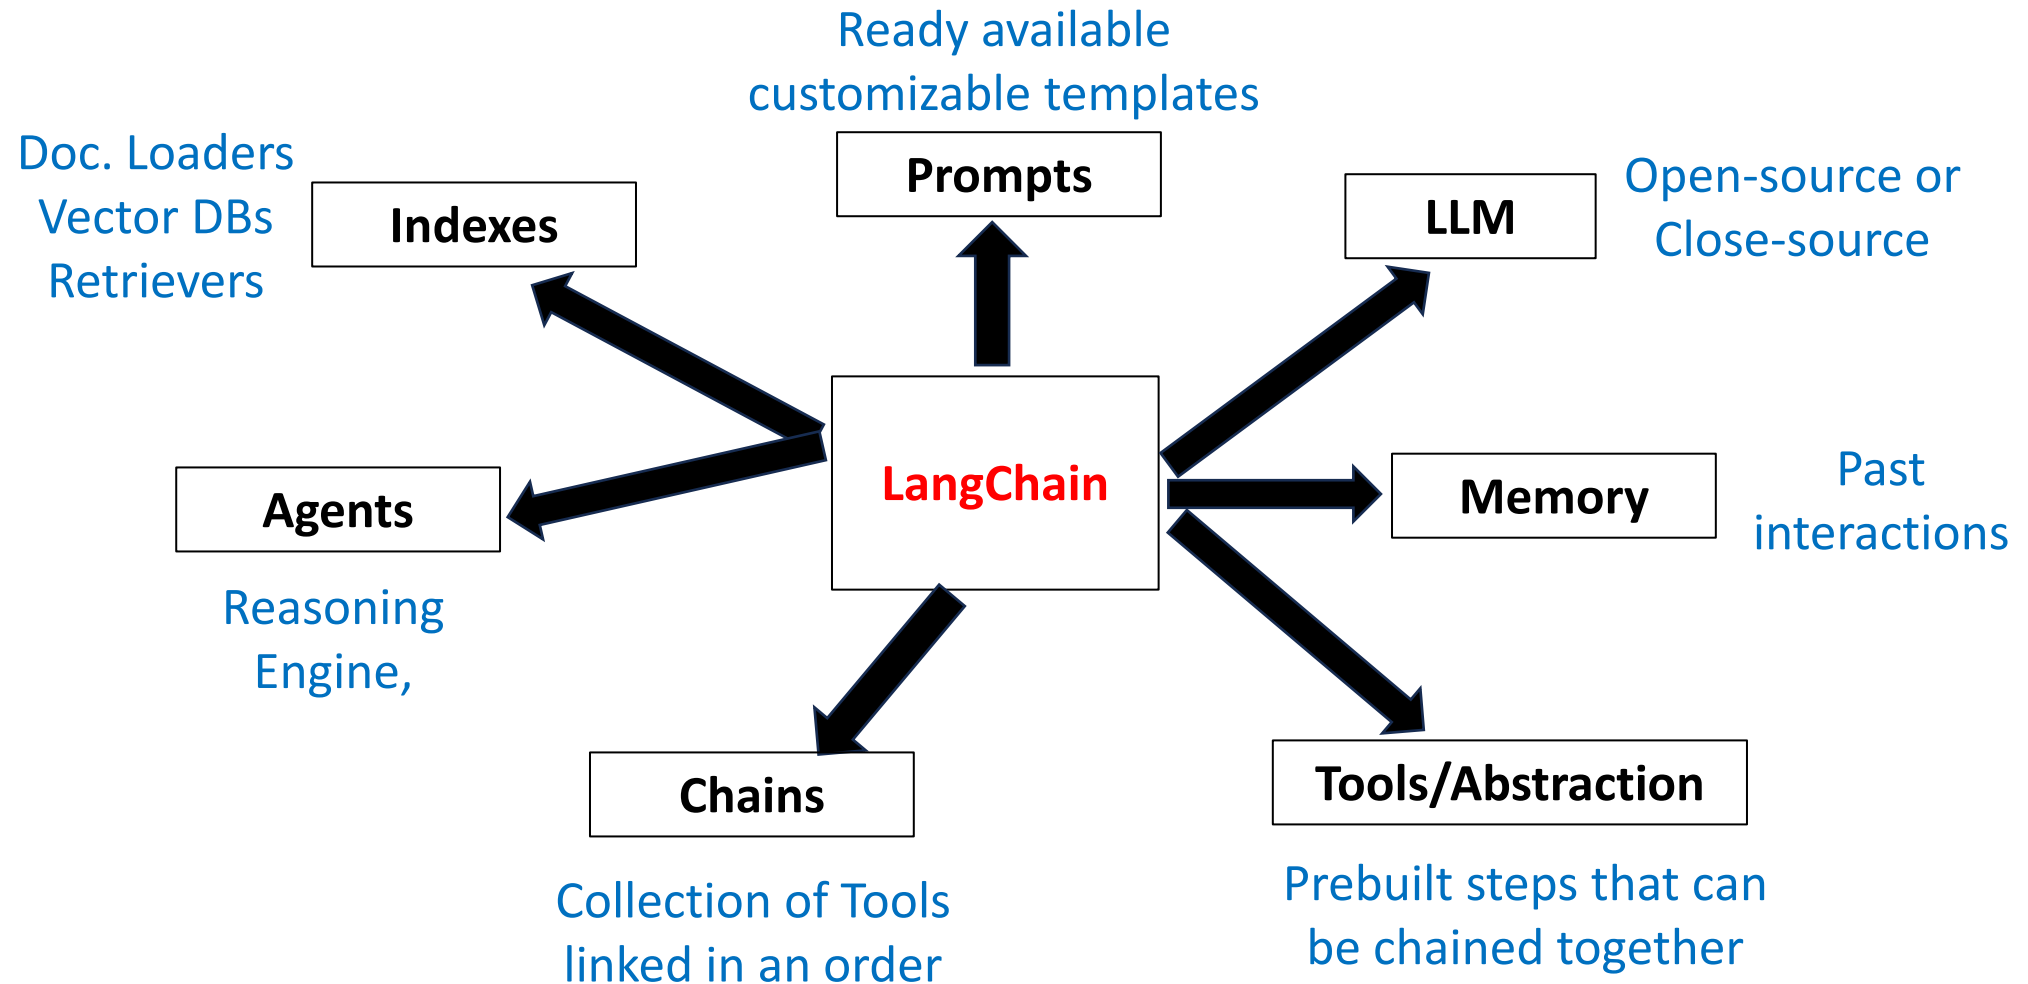
Basic RAG Approach (Recap from Langchain Training)



Langchain can automate the above manual tasks for you

Majorly it can: Interact with External Data, make API Calls, use Memory etc.

LangChain Components



LangChain Components

1. Tools:

1. **Independent components** that can be chained
 - Embeddings, vector stores, Loaders etc.
2. Interfaces that allow language models to **interact with external systems**, such as:
 - APIs
 - Databases
 - Functions etc.
3. Each Tool has:
 - Name
 - Description
 - Inputs
 - Function
4. **Examples:**
 1. A Function that queries a DB
 2. Call an external API
 3. Initiate a document loader or embeddings

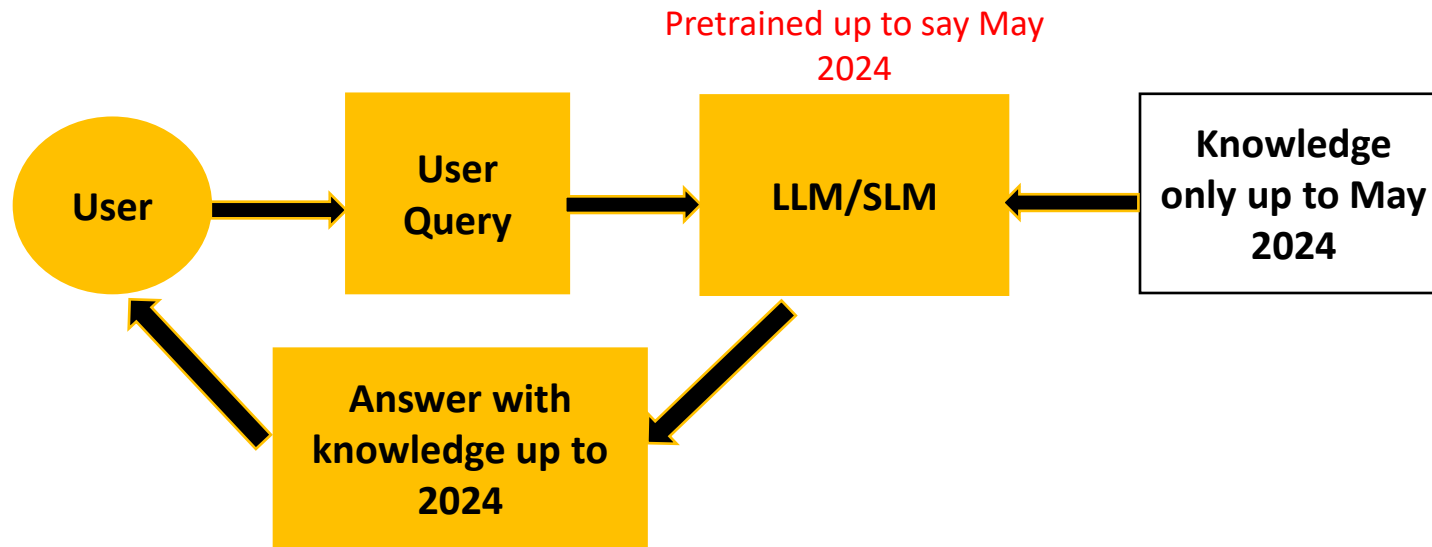
2. Chains:

1. Sequences of Tool calls (or actions): Allows you to **combine multiple tools into a workflow**.
2. Actions in chains are **predefined** and **specified in a specific order**
3. **Linear and Static**. They don't change dynamically based on Input.
4. **Examples:**
 1. **Chain1:** A chain for Data/Doc Processing Pipeline: Loading, Embedding, & Storing
 2. **Chain2:** A chain to Query Handling: combine query embeddings + generate response into a chain

3. Agents:

1. Manage workflows – dynamically decide **which Tool/Chain to use** and in **which order**
2. Agents are responsible for:
 1. Overall logic
 2. Can dynamically adjust the workflow based on input & context
3. **Examples:**
 1. An agent that uses a doc loader, embedding model, vectorstore, & lang model to handle user queries
 2. An agent that process Chain2 after Chain1

Limitation of Pre-trained Models (Recap from MCP Training)



Observations

- Answers are limited to its pretrained ability
- Not open to live or real time information
- For real time info, we need to connect to a tool/agent that can fetch needed information in real time